

FOOD AI & BUSINESS AUTOMATION

WALMART-SCALE RETAIL PATTERNS

A vendor-neutral field guide for grocery, fresh food, retail operations, supply chain, AI agents and enterprise automation.

CUSTOMER

STORE

SUPPLY

DATA

AI

WORKFLOW

GOVERN

104 best practices | 13 operating domains | 100+ pages

Designed for: AI automation architects, retail leaders, engineers, data/ML teams and consultants

How to Use This Handbook

A field guide for building capabilities, not a catalog of AI tools.

- 1 Start with the business decision.** Pick a measurable retail decision - order quantity, substitution, markdown, route, task, invoice exception - and define who owns it before choosing a model.
- 2 Build the trusted data path first.** Every practice assumes authoritative product/location/supplier IDs, reconciled inventory, event timestamps, data quality checks and an auditable source trail.
- 3 Automate progressively.** Begin with observe and recommend, then controlled approval, then auto-execution only for low-risk cases with proven quality and rollback.
- 4 Separate intelligence from action.** Forecasts, optimization and LLMs generate decisions or interpretations; deterministic policy services and workflows enforce constraints before systems of record are changed.
- 5 Measure closed-loop value.** Track availability, waste, labor, margin, working capital, service, safety and reliability - not just model accuracy or token usage.
- 6 Use the chapter pages as implementation tickets.** Each best practice gives a target outcome, production pattern, rollout sequence, controls, KPIs and a concrete first deliverable.

Executive Implementation Principles

The common design rules behind a Walmart-scale food automation platform.

- 1 System of record stays authoritative.** AI should not become a shadow ERP/WMS/OMS. Persist business state in governed transactional systems and expose it through stable APIs/events.
- 2 Inventory confidence is as important as inventory quantity.** Bad stock records poison forecasting, ATP, replenishment, substitutions and customer promises; detect and route low-confidence records explicitly.
- 3 Fresh food is a time-decaying asset.** Remaining shelf life, temperature and lot genealogy must be first-class decision inputs across receiving, allocation, markdown and fulfillment.
- 4 LLMs are best at ambiguity, not transaction integrity.** Use them for unstructured documents, SOP retrieval, explanations and exception triage; use typed services for prices, orders, payments and inventory writes.
- 5 Human control follows consequence.** Safety holds, payment/bank changes, high-value orders, employment impacts and major customer remedies require stronger approvals and evidence.
- 6 Every automation needs a degraded mode.** Stores, warehouses and customer channels must continue operating when a model, agent, third party or network dependency is down.

Target Capability Map

A practical capability stack for a large grocery and food retailer.

1

Commercial: forecasting, pricing, promotions, assortment and space optimization tied to inventory and margin constraints.

2

Supply chain: supplier collaboration, purchase orders, inbound visibility, DC optimization, routing, capacity and control-tower exceptions.

3

Store/fresh: shelf availability, task orchestration, shelf life, markdown, prep planning, waste, cold chain and asset monitoring.

4

Omnichannel: inventory promises, order routing, picking, substitution, delivery promises, customer service and personalization.

5

Safety/traceability: lot genealogy, critical tracking events, evidence, quarantine, recall impact analysis and regulatory readiness.

6

Platform: event/data products, APIs, workflows, agents/RAG, security, evaluation, observability, FinOps and reliable production operations.

Reference Architecture: From Retail Events to Governed Actions

The architecture separates systems of record, data products, intelligence, orchestration and controlled action.

Experience & Operations	Store POS eCommerce associate app supplier portal control tower contact center
Action & Workflow	OMS/WMS/TMS/ERP/CMMS adapters approval gates task queues notifications RPA where no API exists
Decision Services	Forecasting replenishment pricing routing substitution anomaly detection computer vision optimization
Agent & Knowledge	RAG policy/SOP retrieval tool gateway agent runtime conversation state evaluation trace store
Data Products	Sales inventory ledger product/location/supplier master traceability events customer consent feature store
Platform & Governance	Identity secrets event bus API gateway observability model registry CI/CD cost controls audit

Core design rule: AI proposes or interprets; deterministic services validate constraints; workflow orchestration records state; systems of record remain authoritative; humans approve consequential exceptions.

Table of Contents (1/3)

Page numbers refer to the generated handbook structure.

Chapter 01: Strategy, Value Streams & Operating Model.....	9
01.01 Prioritize automation by value, feasibility and risk	10
01.02 Map grocery value streams end to end	11
01.03 Design a human plus AI operating model	12
01.04 Establish product ownership for AI capabilities	13
01.05 Build a benefits realization and KPI framework	14
01.06 Create decision rights and escalation paths	15
01.07 Use process mining before workflow redesign	16
01.08 Build the 12-month transformation portfolio	17
Chapter 02: Retail Data Foundation & Master Data.....	18
02.01 Design a canonical retail data model	19
02.02 Master products, locations, suppliers and customers	20
02.03 Create real-time sales and inventory event streams	21
02.04 Implement data quality contracts and observability	22
02.05 Build a feature platform for retail machine learning	23
02.06 Manage identity, consent and customer data boundaries	24
02.07 Design the semantic layer and metric definitions	25
02.08 Create lineage, retention and data access controls	26
Chapter 03: Demand Forecasting & Scenario Intelligence.....	27
03.01 Forecast demand at SKU-location-day grain	28
03.02 Use hierarchical and reconciled forecasting	29
03.03 Model promotions, holidays, weather and events	30
03.04 Produce probabilistic forecasts and confidence bands	31
03.05 Detect demand regime changes and model drift	32
03.06 Forecast new items and sparse demand	33
03.07 Run scenario planning for shocks and constraints	34
03.08 Operationalize forecast accuracy and bias management	35
Chapter 04: Inventory, Replenishment & Shelf Availability.....	36
04.01 Calculate a trusted real-time inventory position	37
04.02 Automate reorder points and safety stock	38
04.03 Optimize store and DC replenishment orders	39
04.04 Prioritize on-shelf availability exceptions	40
04.05 Detect phantom inventory and stock record errors	41
04.06 Balance service level, waste and working capital	42
04.07 Automate inter-store and network inventory rebalancing	43
04.08 Build replenishment exception management with human approval	44
Chapter 05: Fresh Food, Shelf Life & Waste Automation.....	45
05.01 Capture lot, batch, expiry and shelf-life attributes	46
05.02 Use FEFO inventory allocation and picking	47
05.03 Forecast fresh demand with intra-day granularity	48
05.04 Optimize dynamic markdowns for expiring food	49
05.05 Predict shrink, spoilage and unsellable inventory	50
05.06 Automate fresh production and prep planning	51
05.07 Route surplus to donation, secondary sale or disposal	52

Table of Contents (2/3)

Page numbers refer to the generated handbook structure.

05.08 Monitor cold-chain risk and product quality	53
Chapter 06: Supplier, Procurement & Accounts Payable Automation.....	54
06.01 Create supplier scorecards for OTIF, quality and cost	55
06.02 Automate purchase order creation and approvals	56
06.03 Process supplier ASNs and receiving exceptions	57
06.04 Use three-way match for invoice automation	58
06.05 Detect duplicate, anomalous and high-risk invoices	59
06.06 Automate supplier onboarding and document validation	60
06.07 Predict supplier risk and disruption exposure	61
06.08 Create an AI procurement copilot with bounded actions	62
Chapter 07: Distribution Centers, Transportation & Network Automation.....	63
07.01 Integrate WMS, TMS, OMS and inventory events	64
07.02 Optimize slotting and product placement in DCs	65
07.03 Automate wave, pick, pack and replenishment priorities	66
07.04 Optimize truck loading and cube utilization	67
07.05 Use route optimization with delivery constraints	68
07.06 Automate yard, dock and appointment scheduling	69
07.07 Predict equipment failures and maintenance needs	70
07.08 Create a network control tower for exceptions	71
Chapter 08: Omnichannel Grocery, Fulfillment & Customer AI.....	72
08.01 Build available-to-promise inventory for omnichannel	73
08.02 Optimize digital order routing and fulfillment nodes	74
08.03 Improve picker routing and batch picking	75
08.04 Automate intelligent item substitution	76
08.05 Predict pickup and delivery promise times	77
08.06 Build a governed customer service AI assistant	78
08.07 Personalize recommendations with privacy controls	79
08.08 Unify customer feedback and service recovery workflows	80
Chapter 09: Pricing, Promotions, Assortment & Merchandising Intelligence.....	81
09.01 Estimate price elasticity with causal safeguards	82
09.02 Optimize promotions for incremental profit	83
09.03 Prevent promotion stockouts with integrated planning	84
09.04 Automate competitive price intelligence responsibly	85
09.05 Optimize assortment by store and customer mission	86
09.06 Use space and planogram analytics	87
09.07 Detect cannibalization and halo effects	88
09.08 Create merchant decision support with explainable recommendations	89
Chapter 10: Store Operations, Workforce & Physical Asset Automation.....	90
10.01 Create a unified store task orchestration layer	91
10.02 Use computer vision for shelf and display exceptions	92
10.03 Automate temperature and refrigeration monitoring	93
10.04 Forecast workload and build labor demand curves	94
10.05 Optimize schedules with policy and fairness constraints	95
10.06 Predict asset failures and service needs	96

Table of Contents (3/3)

Page numbers refer to the generated handbook structure.

10.07 Automate opening, closing and compliance checklists	97
10.08 Build an associate copilot for procedures and exceptions	98
Chapter 11: Food Safety, Traceability, Recall & Compliance	99
11.01 Map critical tracking events and key data elements	100
11.02 Implement interoperable lot-level traceability	101
11.03 Use GS1 identifiers and EPCIS-style event records	102
11.04 Automate supplier food-safety document controls	103
11.05 Monitor HACCP and critical control point exceptions	104
11.06 Build a recall impact graph and containment workflow	105
11.07 Preserve audit evidence and chain of custody	106
11.08 Test recall readiness with simulation exercises	107
Chapter 12: AI Agents, Workflow Automation & Integration Architecture	108
12.01 Choose workflow automation versus AI agent patterns	109
12.02 Design tool contracts for agent actions	110
12.03 Use RAG for policy, SOP and product knowledge	111
12.04 Orchestrate multi-agent work with a supervisor pattern	112
12.05 Add approval gates for financial and operational actions	113
12.06 Implement retries, idempotency and dead-letter handling	114
12.07 Use n8n-style orchestration for cross-system workflows	115
12.08 Build agent evaluation, tracing and prompt lifecycle management	116
Chapter 13: Platform Engineering, Security, Governance & FinOps	117
13.01 Create an AI platform reference architecture	118
13.02 Apply least privilege, secrets management and network controls	119
13.03 Classify AI use cases by risk and required controls	120
13.04 Implement model, prompt and workflow versioning	121
13.05 Monitor latency, errors, quality and business outcomes	122
13.06 Control LLM and automation cost with budgets and routing	123
13.07 Design disaster recovery and graceful degradation	124
13.08 Run production release, rollback and incident management	125
Appendix A: Canonical Retail Data Model	126
Appendix B: Example Retail Event Contract	127
Appendix C: Replenishment Decision Pseudocode	128
Appendix D: Governed Retail Agent Pattern	129
Appendix E: Retail AI KPI Dictionary	130
Appendix F: 12-Month Implementation Roadmap	131
Appendix G: Team & RACI Blueprint	132
Appendix H: Vendor / Platform Evaluation Scorecard	133
Appendix I: AI / Automation Risk Classification	134
Appendix J: Production Go-Live Checklist	135
Appendix K: Incident & Rollback Playbook	136
Appendix L: Food Traceability Implementation Note	137
Appendix M: 30-60-90 Day Starter Plan	138
References & Further Reading	139

CHAPTER 01

Strategy, Value Streams & Operating Model

Define the business outcomes, operating model, decision rights, and value measurement before automating work.

What good looks like

- One accountable business owner owns the decision and its KPI; technology ownership is explicit but does not replace operational accountability.
- Inputs are reconciled to systems of record with freshness and quality SLOs; recommendations carry confidence, timestamp, reason codes and version lineage.
- Automation is progressive: observe -> recommend -> approve -> auto-execute low-risk cases. High-impact exceptions retain human control and audit evidence.
- The capability is evaluated on business outcomes and operational reliability, not model accuracy alone. A fallback mode keeps stores and supply chain running.

Best practices in this chapter

01.01 Prioritize automation by value, feasibility and risk

01.05 Build a benefits realization and KPI framework

01.02 Map grocery value streams end to end

01.06 Create decision rights and escalation paths

01.03 Design a human plus AI operating model

01.07 Use process mining before workflow redesign

01.04 Establish product ownership for AI capabilities

01.08 Build the 12-month transformation portfolio

- Authoritative inputs: business case, process baseline, ownership and change management.
- Expected implementation assets: value stream map, KPI tree, RACI, risk register and benefit baseline.
- Required cross-cutting controls: identity, audit, data quality, observability, rollback and human escalation.

MINIMUM CAPABILITY FOUNDATION

CHAPTER 01 / PRACTICE 01

Prioritize automation by value, feasibility and risk

Strategy, Value Streams & Operating Model

- Implement **prioritize automation by value, feasibility and risk** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Classify use cases by impact, autonomy, data sensitivity and reversibility; require stronger evidence and human control as consequence severity rises. Use business case, process baseline, ownership and change management. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from business case, process baseline, ownership and change management. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- validated annual value
- cycle-time reduction
- user adoption / override rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: value stream map, KPI tree, RACI, risk register and benefit baseline; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 01 / PRACTICE 02

Map grocery value streams end to end


Strategy, Value Streams & Operating Model

- Implement **map grocery value streams end to end** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use business case, process baseline, ownership and change management. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from business case, process baseline, ownership and change management. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- validated annual value
- cycle-time reduction
- user adoption / override rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: value stream map, KPI tree, RACI, risk register and benefit baseline; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 01 / PRACTICE 03

Design a human plus AI operating model


Strategy, Value Streams & Operating Model

- Implement **design a human plus ai operating model** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use business case, process baseline, ownership and change management. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from business case, process baseline, ownership and change management. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- validated annual value
- cycle-time reduction
- user adoption / override rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: value stream map, KPI tree, RACI, risk register and benefit baseline; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 01 / PRACTICE 04

Establish product ownership for AI capabilities

Strategy, Value Streams & Operating Model

- Implement **establish product ownership for ai capabilities** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use business case, process baseline, ownership and change management. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from business case, process baseline, ownership and change management. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- validated annual value
- cycle-time reduction
- user adoption / override rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: value stream map, KPI tree, RACI, risk register and benefit baseline; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 01 / PRACTICE 05

Build a benefits realization and KPI framework


Strategy, Value Streams & Operating Model

- Implement **build a benefits realization and kpi framework** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use business case, process baseline, ownership and change management. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from business case, process baseline, ownership and change management. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- validated annual value
- cycle-time reduction
- user adoption / override rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: value stream map, KPI tree, RACI, risk register and benefit baseline; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 01 / PRACTICE 06

Create decision rights and escalation paths


Strategy, Value Streams & Operating Model

- Implement **create decision rights and escalation paths** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use business case, process baseline, ownership and change management. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from business case, process baseline, ownership and change management. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- validated annual value
- cycle-time reduction
- user adoption / override rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: value stream map, KPI tree, RACI, risk register and benefit baseline; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 01 / PRACTICE 07

Use process mining before workflow redesign


Strategy, Value Streams & Operating Model

- Implement **use process mining before workflow redesign** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use business case, process baseline, ownership and change management. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from business case, process baseline, ownership and change management. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- validated annual value
- cycle-time reduction
- user adoption / override rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: value stream map, KPI tree, RACI, risk register and benefit baseline; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 01 / PRACTICE 08

Build the 12-month transformation portfolio


Strategy, Value Streams & Operating Model

- Implement **build the 12-month transformation portfolio** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use business case, process baseline, ownership and change management. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from business case, process baseline, ownership and change management. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- validated annual value
- cycle-time reduction
- user adoption / override rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: value stream map, KPI tree, RACI, risk register and benefit baseline; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 02

Retail Data Foundation & Master Data

Create trusted, event-driven retail data that can support forecasting, optimization, agents and auditability.

What good looks like

- One accountable business owner owns the decision and its KPI; technology ownership is explicit but does not replace operational accountability.
- Inputs are reconciled to systems of record with freshness and quality SLOs; recommendations carry confidence, timestamp, reason codes and version lineage.
- Automation is progressive: observe -> recommend -> approve -> auto-execute low-risk cases. High-impact exceptions retain human control and audit evidence.
- The capability is evaluated on business outcomes and operational reliability, not model accuracy alone. A fallback mode keeps stores and supply chain running.

Best practices in this chapter

02.01 Design a canonical retail data model

02.05 Build a feature platform for retail machine learning

02.02 Master products, locations, suppliers and customers

02.06 Manage identity, consent and customer data boundaries

02.03 Create real-time sales and inventory event streams

02.07 Design the semantic layer and metric definitions

02.04 Implement data quality contracts and observability

02.08 Create lineage, retention and data access controls

- Authoritative inputs: POS, OMS, WMS, ERP, product, location, supplier and customer data.
- Expected implementation assets: canonical schemas, data contracts, quality SLOs, lineage and access policies.
- Required cross-cutting controls: identity, audit, data quality, observability, rollback and human escalation.

MINIMUM CAPABILITY FOUNDATION

CHAPTER 02 / PRACTICE 01

Design a canonical retail data model


Retail Data Foundation & Master Data

- Implement **design a canonical retail data model** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use POS, OMS, WMS, ERP, product, location, supplier and customer data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from POS, OMS, WMS, ERP, product, location, supplier and customer data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- data freshness SLO
- record reconciliation accuracy
- critical data-quality incident rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: canonical schemas, data contracts, quality SLOs, lineage and access policies; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 02 / PRACTICE 02

Master products, locations, suppliers and customers


Retail Data Foundation & Master Data

- Implement **master products, locations, suppliers and customers** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use POS, OMS, WMS, ERP, product, location, supplier and customer data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from POS, OMS, WMS, ERP, product, location, supplier and customer data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- data freshness SLO
- record reconciliation accuracy
- critical data-quality incident rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: canonical schemas, data contracts, quality SLOs, lineage and access policies; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 02 / PRACTICE 03

Create real-time sales and inventory event streams


Retail Data Foundation & Master Data

- Implement **create real-time sales and inventory event streams** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use POS, OMS, WMS, ERP, product, location, supplier and customer data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from POS, OMS, WMS, ERP, product, location, supplier and customer data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- data freshness SLO
- record reconciliation accuracy
- critical data-quality incident rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: canonical schemas, data contracts, quality SLOs, lineage and access policies; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 02 / PRACTICE 04

Implement data quality contracts and observability


Retail Data Foundation & Master Data

- Implement **implement data quality contracts and observability** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use POS, OMS, WMS, ERP, product, location, supplier and customer data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from POS, OMS, WMS, ERP, product, location, supplier and customer data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- data freshness SLO
- record reconciliation accuracy
- critical data-quality incident rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: canonical schemas, data contracts, quality SLOs, lineage and access policies; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 02 / PRACTICE 05

Build a feature platform for retail machine learning


Retail Data Foundation & Master Data

- Implement **build a feature platform for retail machine learning** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use POS, OMS, WMS, ERP, product, location, supplier and customer data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from POS, OMS, WMS, ERP, product, location, supplier and customer data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- data freshness SLO
- record reconciliation accuracy
- critical data-quality incident rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: canonical schemas, data contracts, quality SLOs, lineage and access policies; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 02 / PRACTICE 06

Manage identity, consent and customer data boundaries


Retail Data Foundation & Master Data

- Implement **manage identity, consent and customer data boundaries** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use POS, OMS, WMS, ERP, product, location, supplier and customer data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from POS, OMS, WMS, ERP, product, location, supplier and customer data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- data freshness SLO
- record reconciliation accuracy
- critical data-quality incident rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: canonical schemas, data contracts, quality SLOs, lineage and access policies; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 02 / PRACTICE 07

Design the semantic layer and metric definitions


Retail Data Foundation & Master Data

- Implement **design the semantic layer and metric definitions** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use POS, OMS, WMS, ERP, product, location, supplier and customer data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from POS, OMS, WMS, ERP, product, location, supplier and customer data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- data freshness SLO
- record reconciliation accuracy
- critical data-quality incident rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: canonical schemas, data contracts, quality SLOs, lineage and access policies; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 02 / PRACTICE 08

Create lineage, retention and data access controls


Retail Data Foundation & Master Data

- Implement **create lineage, retention and data access controls** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use POS, OMS, WMS, ERP, product, location, supplier and customer data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from POS, OMS, WMS, ERP, product, location, supplier and customer data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- data freshness SLO
- record reconciliation accuracy
- critical data-quality incident rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: canonical schemas, data contracts, quality SLOs, lineage and access policies; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 03

Demand Forecasting & Scenario Intelligence

Predict demand at the decision grain and expose uncertainty so planners can make economically correct choices.

What good looks like

- One accountable business owner owns the decision and its KPI; technology ownership is explicit but does not replace operational accountability.
- Inputs are reconciled to systems of record with freshness and quality SLOs; recommendations carry confidence, timestamp, reason codes and version lineage.
- Automation is progressive: observe -> recommend -> approve -> auto-execute low-risk cases. High-impact exceptions retain human control and audit evidence.
- The capability is evaluated on business outcomes and operational reliability, not model accuracy alone. A fallback mode keeps stores and supply chain running.

Best practices in this chapter

03.01 Forecast demand at SKU-location-day grain

03.05 Detect demand regime changes and model drift

03.02 Use hierarchical and reconciled forecasting

03.06 Forecast new items and sparse demand

03.03 Model promotions, holidays, weather and events

03.07 Run scenario planning for shocks and constraints

03.04 Produce probabilistic forecasts and confidence bands

03.08 Operationalize forecast accuracy and bias management

- Authoritative inputs: transactions, inventory, price, promotion, calendar, weather and event features.
- Expected implementation assets: forecast service, backtests, confidence intervals, drift alerts and planner views.
- Required cross-cutting controls: identity, audit, data quality, observability, rollback and human escalation.

MINIMUM CAPABILITY FOUNDATION

CHAPTER 03 / PRACTICE 01

Forecast demand at SKU-location-day grain


Demand Forecasting & Scenario Intelligence

- Implement **forecast demand at sku-location-day grain** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use rolling-origin backtests, compare against a naive seasonal baseline, measure weighted absolute error plus bias, and expose prediction intervals rather than a single point estimate. Use transactions, inventory, price, promotion, calendar, weather and event features. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from transactions, inventory, price, promotion, calendar, weather and event features. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- WAPE / weighted error
- forecast bias
- decision value versus baseline

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: forecast service, backtests, confidence intervals, drift alerts and planner views; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 03 / PRACTICE 02

Use hierarchical and reconciled forecasting


Demand Forecasting & Scenario Intelligence

- Implement **use hierarchical and reconciled forecasting** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use rolling-origin backtests, compare against a naive seasonal baseline, measure weighted absolute error plus bias, and expose prediction intervals rather than a single point estimate. Forecast at decision-relevant grains, reconcile SKU/store/category totals, and preserve additivity so finance, merchandising and replenishment see consistent numbers. Use transactions, inventory, price, promotion, calendar, weather and event features. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from transactions, inventory, price, promotion, calendar, weather and event features. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- WAPE / weighted error
- forecast bias
- decision value versus baseline

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: forecast service, backtests, confidence intervals, drift alerts and planner views; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 03 / PRACTICE 03

Model promotions, holidays, weather and events


Demand Forecasting & Scenario Intelligence

- Implement **model promotions, holidays, weather and events** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use weather and external signals only where uplift is demonstrable; version feature sources, monitor availability, and fall back safely when third-party data is missing. Estimate incrementality, cannibalization, halo and inventory availability; optimize profit or strategic objective rather than raw sales uplift. Use transactions, inventory, price, promotion, calendar, weather and event features. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from transactions, inventory, price, promotion, calendar, weather and event features. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- WAPE / weighted error
- forecast bias
- decision value versus baseline

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: forecast service, backtests, confidence intervals, drift alerts and planner views; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 03 / PRACTICE 04

Produce probabilistic forecasts and confidence bands


Demand Forecasting & Scenario Intelligence

- Implement **produce probabilistic forecasts and confidence bands** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use rolling-origin backtests, compare against a naive seasonal baseline, measure weighted absolute error plus bias, and expose prediction intervals rather than a single point estimate. Return quantiles such as P50/P80/P95 and translate uncertainty into safety stock, capacity and staffing decisions instead of hiding it in a single forecast. Use transactions, inventory, price, promotion, calendar, weather and event features. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from transactions, inventory, price, promotion, calendar, weather and event features. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- WAPE / weighted error
- forecast bias
- decision value versus baseline

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: forecast service, backtests, confidence intervals, drift alerts and planner views; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 03 / PRACTICE 05

Detect demand regime changes and model drift


Demand Forecasting & Scenario Intelligence

- Implement **detect demand regime changes and model drift** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use transactions, inventory, price, promotion, calendar, weather and event features. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from transactions, inventory, price, promotion, calendar, weather and event features. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- WAPE / weighted error
- forecast bias
- decision value versus baseline

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: forecast service, backtests, confidence intervals, drift alerts and planner views; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 03 / PRACTICE 06

Forecast new items and sparse demand

Demand Forecasting & Scenario Intelligence

- Implement **forecast new items and sparse demand** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use rolling-origin backtests, compare against a naive seasonal baseline, measure weighted absolute error plus bias, and expose prediction intervals rather than a single point estimate. Use attribute similarity, category priors and launch curves; require planner overrides for major launches and rapidly re-weight toward observed demand after launch. Use transactions, inventory, price, promotion, calendar, weather and event features. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from transactions, inventory, price, promotion, calendar, weather and event features. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- WAPE / weighted error
- forecast bias
- decision value versus baseline

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: forecast service, backtests, confidence intervals, drift alerts and planner views; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 03 / PRACTICE 07

Run scenario planning for shocks and constraints

Demand Forecasting & Scenario Intelligence

- Implement **run scenario planning for shocks and constraints** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use transactions, inventory, price, promotion, calendar, weather and event features. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from transactions, inventory, price, promotion, calendar, weather and event features. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- WAPE / weighted error
- forecast bias
- decision value versus baseline

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: forecast service, backtests, confidence intervals, drift alerts and planner views; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 03 / PRACTICE 08

Operationalize forecast accuracy and bias management


Demand Forecasting & Scenario Intelligence

- Implement **operationalize forecast accuracy and bias management** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use rolling-origin backtests, compare against a naive seasonal baseline, measure weighted absolute error plus bias, and expose prediction intervals rather than a single point estimate. Use transactions, inventory, price, promotion, calendar, weather and event features. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from transactions, inventory, price, promotion, calendar, weather and event features. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- WAPE / weighted error
- forecast bias
- decision value versus baseline

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: forecast service, backtests, confidence intervals, drift alerts and planner views; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 04

Inventory, Replenishment & Shelf Availability

Turn forecasts and inventory signals into constrained ordering decisions that protect availability and working capital.

What good looks like

- One accountable business owner owns the decision and its KPI; technology ownership is explicit but does not replace operational accountability.
- Inputs are reconciled to systems of record with freshness and quality SLOs; recommendations carry confidence, timestamp, reason codes and version lineage.
- Automation is progressive: observe -> recommend -> approve -> auto-execute low-risk cases. High-impact exceptions retain human control and audit evidence.
- The capability is evaluated on business outcomes and operational reliability, not model accuracy alone. A fallback mode keeps stores and supply chain running.

Best practices in this chapter

04.01 Calculate a trusted real-time inventory position

04.05 Detect phantom inventory and stock record errors

04.02 Automate reorder points and safety stock

04.06 Balance service level, waste and working capital

04.03 Optimize store and DC replenishment orders

04.07 Automate inter-store and network inventory rebalancing

04.04 Prioritize on-shelf availability exceptions

04.08 Build replenishment exception management with human approval

- Authoritative inputs: inventory position, lead time, order cycles, pack sizes, service targets and constraints.
- Expected implementation assets: inventory ledger, replenishment policy, optimizer, exception queue and approvals.
- Required cross-cutting controls: identity, audit, data quality, observability, rollback and human escalation.

MINIMUM CAPABILITY FOUNDATION

CHAPTER 04 / PRACTICE 01

Calculate a trusted real-time inventory position


Inventory, Replenishment & Shelf Availability

- Implement **calculate a trusted real-time inventory position** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Maintain an event-sourced ledger for on-hand, reserved, in-transit, damaged, quarantine and expected receipts; reconcile physical counts and surface confidence in the balance. Use inventory position, lead time, order cycles, pack sizes, service targets and constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from inventory position, lead time, order cycles, pack sizes, service targets and constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- on-shelf availability
- stockout rate
- inventory days / turns

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: inventory ledger, replenishment policy, optimizer, exception queue and approvals; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 04 / PRACTICE 02

Automate reorder points and safety stock

Inventory, Replenishment & Shelf Availability

- Implement **automate reorder points and safety stock** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Calculate service-level targets from demand and lead-time uncertainty; constrain by shelf life, case pack, minimum order quantity and storage capacity. Use inventory position, lead time, order cycles, pack sizes, service targets and constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from inventory position, lead time, order cycles, pack sizes, service targets and constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- on-shelf availability
- stockout rate
- inventory days / turns

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: inventory ledger, replenishment policy, optimizer, exception queue and approvals; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 04 / PRACTICE 03

Optimize store and DC replenishment orders


Inventory, Replenishment & Shelf Availability

- Implement **optimize store and dc replenishment orders** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Separate policy from execution: compute recommended order quantities, explain the drivers, apply hard constraints, then auto-release only low-risk orders inside approved thresholds. Use inventory position, lead time, order cycles, pack sizes, service targets and constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from inventory position, lead time, order cycles, pack sizes, service targets and constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- on-shelf availability
- stockout rate
- inventory days / turns

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: inventory ledger, replenishment policy, optimizer, exception queue and approvals; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 04 / PRACTICE 04

Prioritize on-shelf availability exceptions

Inventory, Replenishment & Shelf Availability

- Implement **prioritize on-shelf availability exceptions** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use inventory position, lead time, order cycles, pack sizes, service targets and constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from inventory position, lead time, order cycles, pack sizes, service targets and constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- on-shelf availability
- stockout rate
- inventory days / turns

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: inventory ledger, replenishment policy, optimizer, exception queue and approvals; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 04 / PRACTICE 05

Detect phantom inventory and stock record errors


Inventory, Replenishment & Shelf Availability

- Implement **detect phantom inventory and stock record errors** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Detect mismatches using sales velocity, shelf observations, receiving history and repeated zero-pick signals; trigger cycle counts before changing authoritative inventory. Use inventory position, lead time, order cycles, pack sizes, service targets and constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from inventory position, lead time, order cycles, pack sizes, service targets and constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- on-shelf availability
- stockout rate
- inventory days / turns

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: inventory ledger, replenishment policy, optimizer, exception queue and approvals; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 04 / PRACTICE 06

Balance service level, waste and working capital

Inventory, Replenishment & Shelf Availability

- Implement **balance service level, waste and working capital** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use inventory position, lead time, order cycles, pack sizes, service targets and constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from inventory position, lead time, order cycles, pack sizes, service targets and constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- on-shelf availability
- stockout rate
- inventory days / turns

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: inventory ledger, replenishment policy, optimizer, exception queue and approvals; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 04 / PRACTICE 07

Automate inter-store and network inventory rebalancing

Inventory, Replenishment & Shelf Availability

- Implement **automate inter-store and network inventory rebalancing** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Optimize transfers only after accounting for transfer labor, transport cost, remaining shelf life and source-store service risk; prevent transferring bad inventory records. Use inventory position, lead time, order cycles, pack sizes, service targets and constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from inventory position, lead time, order cycles, pack sizes, service targets and constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

IMPLEMENTATION SEQUENCE

- on-shelf availability
- stockout rate
- inventory days / turns

KPIS / ACCEPTANCE GATES

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: inventory ledger, replenishment policy, optimizer, exception queue and approvals; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 04 / PRACTICE 08

Build replenishment exception management with human approval


Inventory, Replenishment & Shelf Availability

- Implement **build replenishment exception management with human approval** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Separate policy from execution: compute recommended order quantities, explain the drivers, apply hard constraints, then auto-release only low-risk orders inside approved thresholds. Define thresholds for money, safety, inventory write-offs, supplier changes and customer compensation; require explicit approval tokens for irreversible actions. Use inventory position, lead time, order cycles, pack sizes, service targets and constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from inventory position, lead time, order cycles, pack sizes, service targets and constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- on-shelf availability
- stockout rate
- inventory days / turns

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: inventory ledger, replenishment policy, optimizer, exception queue and approvals; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 05

Fresh Food, Shelf Life & Waste Automation

Manage perishability explicitly with shelf-life-aware inventory, markdown, production and donation decisions.
What good looks like

- One accountable business owner owns the decision and its KPI; technology ownership is explicit but does not replace operational accountability.
- Inputs are reconciled to systems of record with freshness and quality SLOs; recommendations carry confidence, timestamp, reason codes and version lineage.
- Automation is progressive: observe -> recommend -> approve -> auto-execute low-risk cases. High-impact exceptions retain human control and audit evidence.
- The capability is evaluated on business outcomes and operational reliability, not model accuracy alone. A fallback mode keeps stores and supply chain running.

Best practices in this chapter

05.01 Capture lot, batch, expiry and shelf-life attributes

05.02 Use FEFO inventory allocation and picking

05.03 Forecast fresh demand with intra-day granularity

05.04 Optimize dynamic markdowns for expiring food

05.05 Predict shrink, spoilage and unsellable inventory

05.06 Automate fresh production and prep planning

05.07 Route surplus to donation, secondary sale or disposal

05.08 Monitor cold-chain risk and product quality

- Authoritative inputs: lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data.
- Expected implementation assets: FEFO rules, perishability models, markdown engine, waste workflows and quality alerts.
- Required cross-cutting controls: identity, audit, data quality, observability, rollback and human escalation.

MINIMUM CAPABILITY FOUNDATION

CHAPTER 05 / PRACTICE 01

Capture lot, batch, expiry and shelf-life attributes


Fresh Food, Shelf Life & Waste Automation

- Implement **capture lot, batch, expiry and shelf-life attributes** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Store manufacture/pack/expiry dates and remaining shelf life as first-class attributes; reject receipts below agreed life thresholds and preserve lot lineage. Use lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fresh shrink %
- sell-through before expiry
- food waste units/value

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: FEFO rules, perishability models, markdown engine, waste workflows and quality alerts; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 05 / PRACTICE 02

Use FEFO inventory allocation and picking


Fresh Food, Shelf Life & Waste Automation

- Implement **use fefo inventory allocation and picking** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Allocate the earliest valid expiry first while respecting quality holds, customer promise, channel rules and minimum remaining shelf life at delivery. Use lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fresh shrink %
- sell-through before expiry
- food waste units/value

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: FEFO rules, perishability models, markdown engine, waste workflows and quality alerts; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 05 / PRACTICE 03

Forecast fresh demand with intra-day granularity

Fresh Food, Shelf Life & Waste Automation

- Implement **forecast fresh demand with intra-day granularity** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use rolling-origin backtests, compare against a naive seasonal baseline, measure weighted absolute error plus bias, and expose prediction intervals rather than a single point estimate. Use lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

IMPLEMENTATION SEQUENCE

- fresh shrink %
- sell-through before expiry
- food waste units/value

KPIS / ACCEPTANCE GATES

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: FEFO rules, perishability models, markdown engine, waste workflows and quality alerts; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 05 / PRACTICE 04

Optimize dynamic markdowns for expiring food


Fresh Food, Shelf Life & Waste Automation

- Implement **optimize dynamic markdowns for expiring food** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Optimize markdown timing and depth against expected sell-through, margin, waste cost and brand constraints; require deterministic price floors and legal guardrails. Use lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fresh shrink %
- sell-through before expiry
- food waste units/value

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: FEFO rules, perishability models, markdown engine, waste workflows and quality alerts; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 05 / PRACTICE 05

Predict shrink, spoilage and unsellable inventory

Fresh Food, Shelf Life & Waste Automation

- Implement **predict shrink, spoilage and unsellable inventory** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Model waste probability by item, store, age, temperature and demand; use the model to prioritize actions, not to automatically write off inventory without evidence. Use lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fresh shrink %
- sell-through before expiry
- food waste units/value

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: FEFO rules, perishability models, markdown engine, waste workflows and quality alerts; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 05 / PRACTICE 06

Automate fresh production and prep planning

Fresh Food, Shelf Life & Waste Automation

- Implement **automate fresh production and prep planning** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

IMPLEMENTATION SEQUENCE

- fresh shrink %
- sell-through before expiry
- food waste units/value

KPIS / ACCEPTANCE GATES

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: FEFO rules, perishability models, markdown engine, waste workflows and quality alerts; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 05 / PRACTICE 07

Route surplus to donation, secondary sale or disposal


Fresh Food, Shelf Life & Waste Automation

- Implement **route surplus to donation, secondary sale or disposal** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fresh shrink %
- sell-through before expiry
- food waste units/value

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: FEFO rules, perishability models, markdown engine, waste workflows and quality alerts; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 05 / PRACTICE 08

Monitor cold-chain risk and product quality


Fresh Food, Shelf Life & Waste Automation

- Implement **monitor cold-chain risk and product quality** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Stream temperature telemetry, detect sustained threshold excursions, correlate to lot exposure, and route questionable product to quarantine until a qualified decision is recorded. Classify use cases by impact, autonomy, data sensitivity and reversibility; require stronger evidence and human control as consequence severity rises. Use lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot, expiry, remaining shelf life, temperature, sales velocity, waste and markdown data. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fresh shrink %
- sell-through before expiry
- food waste units/value

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: FEFO rules, perishability models, markdown engine, waste workflows and quality alerts; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 06

Supplier, Procurement & Accounts Payable Automation

Connect suppliers through reliable digital transactions and automate low-risk purchasing and settlement work.

What good looks like

- One accountable business owner owns the decision and its KPI; technology ownership is explicit but does not replace operational accountability.
- Inputs are reconciled to systems of record with freshness and quality SLOs; recommendations carry confidence, timestamp, reason codes and version lineage.
- Automation is progressive: observe -> recommend -> approve -> auto-execute low-risk cases. High-impact exceptions retain human control and audit evidence.
- The capability is evaluated on business outcomes and operational reliability, not model accuracy alone. A fallback mode keeps stores and supply chain running.

Best practices in this chapter

06.01 Create supplier scorecards for OTIF, quality and cost

06.05 Detect duplicate, anomalous and high-risk invoices

06.02 Automate purchase order creation and approvals

06.06 Automate supplier onboarding and document validation

06.03 Process supplier ASNs and receiving exceptions

06.07 Predict supplier risk and disruption exposure

06.04 Use three-way match for invoice automation

06.08 Create an AI procurement copilot with bounded actions

- Authoritative inputs: supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events.
- Expected implementation assets: supplier portal/integration, matching rules, risk scores, approval workflow and audit trail.
- Required cross-cutting controls: identity, audit, data quality, observability, rollback and human escalation.

MINIMUM CAPABILITY FOUNDATION

CHAPTER 06 / PRACTICE 01

Create supplier scorecards for OTIF, quality and cost


Supplier, Procurement & Accounts Payable Automation

- Implement **create supplier scorecards for otif, quality and cost** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Measure OTIF, fill rate, defects, lead-time variability, claims and cost variance using governed definitions; show confidence and volume context so small suppliers are not misranked. Track cost per completed business outcome, route simple tasks to smaller models, cache stable results, cap context and set per-workflow budgets and alerts. Use supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- supplier OTIF
- touchless invoice rate
- procure-to-pay exception cycle time

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: supplier portal/integration, matching rules, risk scores, approval workflow and audit trail; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 06 / PRACTICE 02

Automate purchase order creation and approvals


Supplier, Procurement & Accounts Payable Automation

- Implement **automate purchase order creation and approvals** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Generate POs from approved replenishment signals, apply spend/quantity tolerances, validate vendor terms, and require human approval outside policy. Define thresholds for money, safety, inventory write-offs, supplier changes and customer compensation; require explicit approval tokens for irreversible actions. Use supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- supplier OTIF
- touchless invoice rate
- procure-to-pay exception cycle time

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: supplier portal/integration, matching rules, risk scores, approval workflow and audit trail; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 06 / PRACTICE 03

Process supplier ASNs and receiving exceptions


Supplier, Procurement & Accounts Payable Automation

- Implement **process supplier asns and receiving exceptions** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Validate shipment identifiers, quantities, lot/expiry and expected arrival before receipt; route mismatches to an exception queue instead of silently correcting them. Use supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- supplier OTIF
- touchless invoice rate
- procure-to-pay exception cycle time

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: supplier portal/integration, matching rules, risk scores, approval workflow and audit trail; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 06 / PRACTICE 04

Use three-way match for invoice automation


Supplier, Procurement & Accounts Payable Automation

- Implement **use three-way match for invoice automation** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Match PO, receipt and invoice at line level with configurable tolerances; auto-pay clean matches and hold exceptions with reason codes and evidence. Use deterministic duplicate checks plus anomaly models; never let an LLM alone authorize payment or modify bank details. Use supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- supplier OTIF
- touchless invoice rate
- procure-to-pay exception cycle time

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: supplier portal/integration, matching rules, risk scores, approval workflow and audit trail; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 06 / PRACTICE 05

Detect duplicate, anomalous and high-risk invoices


Supplier, Procurement & Accounts Payable Automation

- Implement **detect duplicate, anomalous and high-risk invoices** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use deterministic duplicate checks plus anomaly models; never let an LLM alone authorize payment or modify bank details. Classify use cases by impact, autonomy, data sensitivity and reversibility; require stronger evidence and human control as consequence severity rises. Use supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- supplier OTIF
- touchless invoice rate
- procure-to-pay exception cycle time

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: supplier portal/integration, matching rules, risk scores, approval workflow and audit trail; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 06 / PRACTICE 06

Automate supplier onboarding and document validation


Supplier, Procurement & Accounts Payable Automation

- Implement **automate supplier onboarding and document validation** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- supplier OTIF
- touchless invoice rate
- procure-to-pay exception cycle time

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: supplier portal/integration, matching rules, risk scores, approval workflow and audit trail; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 06 / PRACTICE 07

Predict supplier risk and disruption exposure

Supplier, Procurement & Accounts Payable Automation

- Implement **predict supplier risk and disruption exposure** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Combine operational performance, lead-time volatility, concentration and approved external signals; keep risk drivers explainable and define contingency actions. Classify use cases by impact, autonomy, data sensitivity and reversibility; require stronger evidence and human control as consequence severity rises. Use supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- supplier OTIF
- touchless invoice rate
- procure-to-pay exception cycle time

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: supplier portal/integration, matching rules, risk scores, approval workflow and audit trail; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 06 / PRACTICE 08

Create an AI procurement copilot with bounded actions

Supplier, Procurement & Accounts Payable Automation

- Implement **create an ai procurement copilot with bounded actions** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from supplier master, PO, ASN, receipt, invoice, contract, quality and logistics events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- supplier OTIF
- touchless invoice rate
- procure-to-pay exception cycle time

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: supplier portal/integration, matching rules, risk scores, approval workflow and audit trail; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 07

Distribution Centers, Transportation & Network Automation

Use optimization, telemetry and automation to improve throughput, transport cost, freshness and resilience.

What good looks like

- One accountable business owner owns the decision and its KPI; technology ownership is explicit but does not replace operational accountability.
- Inputs are reconciled to systems of record with freshness and quality SLOs; recommendations carry confidence, timestamp, reason codes and version lineage.
- Automation is progressive: observe -> recommend -> approve -> auto-execute low-risk cases. High-impact exceptions retain human control and audit evidence.
- The capability is evaluated on business outcomes and operational reliability, not model accuracy alone. A fallback mode keeps stores and supply chain running.

Best practices in this chapter

07.01 Integrate WMS, TMS, OMS and inventory events

07.05 Use route optimization with delivery constraints

07.02 Optimize slotting and product placement in DCs

07.06 Automate yard, dock and appointment scheduling

07.03 Automate wave, pick, pack and replenishment priorities

07.07 Predict equipment failures and maintenance needs

07.04 Optimize truck loading and cube utilization

07.08 Create a network control tower for exceptions

- Authoritative inputs: WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state.
- Expected implementation assets: optimization services, control tower, event bus, exception workflows and maintenance triggers.
- Required cross-cutting controls: identity, audit, data quality, observability, rollback and human escalation.

MINIMUM CAPABILITY FOUNDATION

CHAPTER 07 / PRACTICE 01

Integrate WMS, TMS, OMS and inventory events


Distribution Centers, Transportation & Network Automation

- Implement **integrate wms, tms, oms and inventory events** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- cost per case/order
- DC throughput / dock-to-stock
- on-time delivery

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: optimization services, control tower, event bus, exception workflows and maintenance triggers; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 07 / PRACTICE 02

Optimize slotting and product placement in DCs

Distribution Centers, Transportation & Network Automation

• Implement **optimize slotting and product placement in dcs** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

• Optimize pick frequency, cube, weight, affinity, temperature zone and ergonomics; test proposed moves in simulation before large physical re-slotting. Use WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- cost per case/order
- DC throughput / dock-to-stock
- on-time delivery

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: optimization services, control tower, event bus, exception workflows and maintenance triggers; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 07 / PRACTICE 03

Automate wave, pick, pack and replenishment priorities

Distribution Centers, Transportation & Network Automation

- Implement **automate wave, pick, pack and replenishment priorities** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Separate policy from execution: compute recommended order quantities, explain the drivers, apply hard constraints, then auto-release only low-risk orders inside approved thresholds. Use WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- cost per case/order
- DC throughput / dock-to-stock
- on-time delivery

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: optimization services, control tower, event bus, exception workflows and maintenance triggers; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 07 / PRACTICE 04

Optimize truck loading and cube utilization


Distribution Centers, Transportation & Network Automation

- Implement **optimize truck loading and cube utilization** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Respect weight, axle, temperature zones, stop sequence, fragility and cube; produce a load plan that can be validated by warehouse execution systems. Use WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- cost per case/order
- DC throughput / dock-to-stock
- on-time delivery

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: optimization services, control tower, event bus, exception workflows and maintenance triggers; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 07 / PRACTICE 05

Use route optimization with delivery constraints


Distribution Centers, Transportation & Network Automation

- Implement **use route optimization with delivery constraints** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Optimize total cost under time windows, vehicle capacity, driver rules and service priorities; continuously re-plan only when the benefit exceeds disruption cost. Use WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- cost per case/order
- DC throughput / dock-to-stock
- on-time delivery

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: optimization services, control tower, event bus, exception workflows and maintenance triggers; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 07 / PRACTICE 06

Automate yard, dock and appointment scheduling

Distribution Centers, Transportation & Network Automation

- Implement **automate yard, dock and appointment scheduling** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use appointment, trailer, door and labor state to prioritize moves; preserve manual override for safety and unplanned carrier conditions. Use WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- cost per case/order
- DC throughput / dock-to-stock
- on-time delivery

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: optimization services, control tower, event bus, exception workflows and maintenance triggers; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 07 / PRACTICE 07

Predict equipment failures and maintenance needs


Distribution Centers, Transportation & Network Automation

- Implement **predict equipment failures and maintenance needs** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use condition and failure-history features to rank maintenance risk, then integrate with CMMS work orders; avoid replacing manufacturer safety inspections. Use WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- cost per case/order
- DC throughput / dock-to-stock
- on-time delivery

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: optimization services, control tower, event bus, exception workflows and maintenance triggers; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 07 / PRACTICE 08

Create a network control tower for exceptions


Distribution Centers, Transportation & Network Automation

- Implement **create a network control tower for exceptions** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from WMS/TMS/OMS, telemetry, capacity, appointments, routes, service windows and equipment state. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- cost per case/order
- DC throughput / dock-to-stock
- on-time delivery

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: optimization services, control tower, event bus, exception workflows and maintenance triggers; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 08

Omnichannel Grocery, Fulfillment & Customer AI

Orchestrate one inventory promise and one service experience across stores, pickup, delivery and digital channels.

What good looks like

- One accountable business owner owns the decision and its KPI; technology ownership is explicit but does not replace operational accountability.
- Inputs are reconciled to systems of record with freshness and quality SLOs; recommendations carry confidence, timestamp, reason codes and version lineage.
- Automation is progressive: observe -> recommend -> approve -> auto-execute low-risk cases. High-impact exceptions retain human control and audit evidence.
- The capability is evaluated on business outcomes and operational reliability, not model accuracy alone. A fallback mode keeps stores and supply chain running.

Best practices in this chapter

08.01 Build available-to-promise inventory for omnichannel

08.05 Predict pickup and delivery promise times

08.02 Optimize digital order routing and fulfillment nodes

08.06 Build a governed customer service AI assistant

08.03 Improve picker routing and batch picking

08.07 Personalize recommendations with privacy controls

08.04 Automate intelligent item substitution

08.08 Unify customer feedback and service recovery workflows

- Authoritative inputs: ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events.
- Expected implementation assets: order orchestration, promise service, picking optimizer, service AI and feedback loop.
- Required cross-cutting controls: identity, audit, data quality, observability, rollback and human escalation.

MINIMUM CAPABILITY FOUNDATION

CHAPTER 08 / PRACTICE 01

Build available-to-promise inventory for omnichannel


Omnichannel Grocery, Fulfillment & Customer AI

- Implement **build available-to-promise inventory for omnichannel** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Calculate ATP from trusted on-hand, reservations, safety buffers, incoming supply and fulfillment capacity; return confidence and prevent overselling uncertain inventory. Use ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fill rate
- substitution acceptance
- order promise accuracy

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: order orchestration, promise service, picking optimizer, service AI and feedback loop; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 08 / PRACTICE 02

Optimize digital order routing and fulfillment nodes


Omnichannel Grocery, Fulfillment & Customer AI

- Implement **optimize digital order routing and fulfillment nodes** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Choose fulfillment node using inventory confidence, pick capacity, distance, promise, cost and substitution risk; make the objective function explicit. Use ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fill rate
- substitution acceptance
- order promise accuracy

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: order orchestration, promise service, picking optimizer, service AI and feedback loop; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 08 / PRACTICE 03

Improve picker routing and batch picking


Omnichannel Grocery, Fulfillment & Customer AI

- Implement **improve picker routing and batch picking** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Batch compatible orders and route picks by aisle/temperature while respecting cut-off times and food-handling separation rules. Use ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fill rate
- substitution acceptance
- order promise accuracy

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: order orchestration, promise service, picking optimizer, service AI and feedback loop; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 08 / PRACTICE 04

Automate intelligent item substitution


Omnichannel Grocery, Fulfillment & Customer AI

- Implement **automate intelligent item substitution** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Rank substitutes by customer preferences, allergen/diet constraints, size, price and availability; never infer sensitive dietary needs beyond explicit customer data. Use ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fill rate
- substitution acceptance
- order promise accuracy

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: order orchestration, promise service, picking optimizer, service AI and feedback loop; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 08 / PRACTICE 05

Predict pickup and delivery promise times


Omnichannel Grocery, Fulfillment & Customer AI

- Implement **predict pickup and delivery promise times** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fill rate
- substitution acceptance
- order promise accuracy

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: order orchestration, promise service, picking optimizer, service AI and feedback loop; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 08 / PRACTICE 06

Build a governed customer service AI assistant


Omnichannel Grocery, Fulfillment & Customer AI

- Implement **build a governed customer service ai assistant** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Ground answers in order state, policies and approved knowledge; redact sensitive data, authenticate before account actions, and escalate disputes or safety issues. Use ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fill rate
- substitution acceptance
- order promise accuracy

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: order orchestration, promise service, picking optimizer, service AI and feedback loop; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 08 / PRACTICE 07

Personalize recommendations with privacy controls

Omnichannel Grocery, Fulfillment & Customer AI

- Implement **personalize recommendations with privacy controls** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use consented behavior and contextual signals, cap repetitive exposure, separate sponsored ranking, and provide opt-out and explanation controls. Use ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fill rate
- substitution acceptance
- order promise accuracy

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: order orchestration, promise service, picking optimizer, service AI and feedback loop; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 08 / PRACTICE 08

Unify customer feedback and service recovery workflows


Omnichannel Grocery, Fulfillment & Customer AI

- Implement **unify customer feedback and service recovery workflows** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from ATP inventory, baskets, digital orders, substitutions, pick paths, customer contact and delivery events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- fill rate
- substitution acceptance
- order promise accuracy

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: order orchestration, promise service, picking optimizer, service AI and feedback loop; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 09

Pricing, Promotions, Assortment & Merchandising Intelligence

Improve commercial decisions with elasticity, incrementality and constraints instead of opaque price recommendations.

What good looks like

- One accountable business owner owns the decision and its KPI; technology ownership is explicit but does not replace operational accountability.
- Inputs are reconciled to systems of record with freshness and quality SLOs; recommendations carry confidence, timestamp, reason codes and version lineage.
- Automation is progressive: observe -> recommend -> approve -> auto-execute low-risk cases. High-impact exceptions retain human control and audit evidence.
- The capability is evaluated on business outcomes and operational reliability, not model accuracy alone. A fallback mode keeps stores and supply chain running.

Best practices in this chapter

09.01 Estimate price elasticity with causal safeguards

09.05 Optimize assortment by store and customer mission

09.02 Optimize promotions for incremental profit

09.06 Use space and planogram analytics

09.03 Prevent promotion stockouts with integrated planning

09.07 Detect cannibalization and halo effects

09.04 Automate competitive price intelligence responsibly

09.08 Create merchant decision support with explainable recommendations

- Authoritative inputs: price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints.
- Expected implementation assets: pricing guardrails, promo optimizer, merchant cockpit, experimentation and explainability.
- Required cross-cutting controls: identity, audit, data quality, observability, rollback and human escalation.

MINIMUM CAPABILITY FOUNDATION

CHAPTER 09 / PRACTICE 01

Estimate price elasticity with causal safeguards

Pricing, Promotions, Assortment & Merchandising Intelligence

- Implement **estimate price elasticity with causal safeguards** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use causal or quasi-experimental methods where possible; separate correlation from response and apply legal, margin and customer-trust guardrails. Use price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- incremental gross profit
- promo forecast error
- price/assortment exception rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: pricing guardrails, promo optimizer, merchant cockpit, experimentation and explainability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 09 / PRACTICE 02

Optimize promotions for incremental profit


Pricing, Promotions, Assortment & Merchandising Intelligence

- Implement **optimize promotions for incremental profit** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Estimate incrementality, cannibalization, halo and inventory availability; optimize profit or strategic objective rather than raw sales uplift. Use price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- incremental gross profit
- promo forecast error
- price/assortment exception rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: pricing guardrails, promo optimizer, merchant cockpit, experimentation and explainability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 09 / PRACTICE 03

Prevent promotion stockouts with integrated planning

Pricing, Promotions, Assortment & Merchandising Intelligence

- Implement **prevent promotion stockouts with integrated planning** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- incremental gross profit
- promo forecast error
- price/assortment exception rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: pricing guardrails, promo optimizer, merchant cockpit, experimentation and explainability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 09 / PRACTICE 04

Automate competitive price intelligence responsibly


Pricing, Promotions, Assortment & Merchandising Intelligence

- Implement **automate competitive price intelligence responsibly** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use permitted data sources, normalize product matches carefully, retain source provenance, and prohibit automated reactions that violate pricing policy or law. Use price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- incremental gross profit
- promo forecast error
- price/assortment exception rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: pricing guardrails, promo optimizer, merchant cockpit, experimentation and explainability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 09 / PRACTICE 05

Optimize assortment by store and customer mission


Pricing, Promotions, Assortment & Merchandising Intelligence

- Implement **optimize assortment by store and customer mission** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Optimize local assortment using demand, substitution, shelf capacity, strategic roles and supplier constraints; protect essential and compliance-sensitive items. Use price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- incremental gross profit
- promo forecast error
- price/assortment exception rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: pricing guardrails, promo optimizer, merchant cockpit, experimentation and explainability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 09 / PRACTICE 06

Use space and planogram analytics

Pricing, Promotions, Assortment & Merchandising Intelligence

- Implement **use space and planogram analytics** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Compare observed shelf state to approved plan, measure confidence, and route uncertain vision detections for confirmation before generating corrective tasks. Use price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- incremental gross profit
- promo forecast error
- price/assortment exception rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: pricing guardrails, promo optimizer, merchant cockpit, experimentation and explainability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 09 / PRACTICE 07

Detect cannibalization and halo effects


Pricing, Promotions, Assortment & Merchandising Intelligence

- Implement **detect cannibalization and halo effects** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- incremental gross profit
- promo forecast error
- price/assortment exception rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: pricing guardrails, promo optimizer, merchant cockpit, experimentation and explainability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 09 / PRACTICE 08

Create merchant decision support with explainable recommendations

Pricing, Promotions, Assortment & Merchandising Intelligence

- Implement **create merchant decision support with explainable recommendations** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use consented behavior and contextual signals, cap repetitive exposure, separate sponsored ranking, and provide opt-out and explanation controls. Use price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from price history, costs, promotions, competitors, elasticity, assortment, space and inventory constraints. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

IMPLEMENTATION SEQUENCE

- incremental gross profit
- promo forecast error
- price/assortment exception rate

KPIS / ACCEPTANCE GATES

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: pricing guardrails, promo optimizer, merchant cockpit, experimentation and explainability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 10

Store Operations, Workforce & Physical Asset Automation

Convert exceptions into prioritized store tasks and improve labor, equipment and execution without over-automating people decisions.

What good looks like

- One accountable business owner owns the decision and its KPI; technology ownership is explicit but does not replace operational accountability.
- Inputs are reconciled to systems of record with freshness and quality SLOs; recommendations carry confidence, timestamp, reason codes and version lineage.
- Automation is progressive: observe -> recommend -> approve -> auto-execute low-risk cases. High-impact exceptions retain human control and audit evidence.
- The capability is evaluated on business outcomes and operational reliability, not model accuracy alone. A fallback mode keeps stores and supply chain running.

Best practices in this chapter

10.01 Create a unified store task orchestration layer

10.05 Optimize schedules with policy and fairness constraints

10.02 Use computer vision for shelf and display exceptions

10.06 Predict asset failures and service needs

10.03 Automate temperature and refrigeration monitoring

10.07 Automate opening, closing and compliance checklists

10.04 Forecast workload and build labor demand curves

10.08 Build an associate copilot for procedures and exceptions

- Authoritative inputs: tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events.
- Expected implementation assets: task engine, vision/IoT signals, workforce rules, asset alerts and associate copilot.
- Required cross-cutting controls: identity, audit, data quality, observability, rollback and human escalation.

MINIMUM CAPABILITY FOUNDATION

CHAPTER 10 / PRACTICE 01

Create a unified store task orchestration layer


Store Operations, Workforce & Physical Asset Automation

- Implement **create a unified store task orchestration layer** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- task completion on time
- labor hours per unit of work
- asset downtime

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: task engine, vision/IoT signals, workforce rules, asset alerts and associate copilot; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 10 / PRACTICE 02

Use computer vision for shelf and display exceptions


Store Operations, Workforce & Physical Asset Automation

- Implement **use computer vision for shelf and display exceptions** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define image retention and privacy boundaries, measure precision/recall by store conditions, and design a human confirmation path for costly or employee-related decisions. Use tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- task completion on time
- labor hours per unit of work
- asset downtime

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: task engine, vision/IoT signals, workforce rules, asset alerts and associate copilot; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 10 / PRACTICE 03

Automate temperature and refrigeration monitoring


Store Operations, Workforce & Physical Asset Automation

- Implement **automate temperature and refrigeration monitoring** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use calibrated sensors, health checks and gap detection; distinguish a sensor failure from a true excursion and retain the evidence required for food-safety review. Use tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- task completion on time
- labor hours per unit of work
- asset downtime

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: task engine, vision/IoT signals, workforce rules, asset alerts and associate copilot; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 10 / PRACTICE 04

Forecast workload and build labor demand curves

Store Operations, Workforce & Physical Asset Automation

- Implement **forecast workload and build labor demand curves** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use rolling-origin backtests, compare against a naive seasonal baseline, measure weighted absolute error plus bias, and expose prediction intervals rather than a single point estimate. Forecast workload from transactions, deliveries, digital orders and tasks; schedule under labor law, contracts, skills, availability, fairness and minimum staffing constraints. Use tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- task completion on time
- labor hours per unit of work
- asset downtime

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: task engine, vision/IoT signals, workforce rules, asset alerts and associate copilot; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 10 / PRACTICE 05

Optimize schedules with policy and fairness constraints

Store Operations, Workforce & Physical Asset Automation

- Implement **optimize schedules with policy and fairness constraints** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Treat optimization as recommendation under explicit policies; provide reason codes and manager override, and audit any automated changes that affect employees. Use tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- task completion on time
- labor hours per unit of work
- asset downtime

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: task engine, vision/IoT signals, workforce rules, asset alerts and associate copilot; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 10 / PRACTICE 06

Predict asset failures and service needs


Store Operations, Workforce & Physical Asset Automation

- Implement **predict asset failures and service needs** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- task completion on time
- labor hours per unit of work
- asset downtime

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: task engine, vision/IoT signals, workforce rules, asset alerts and associate copilot; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 10 / PRACTICE 07

Automate opening, closing and compliance checklists


Store Operations, Workforce & Physical Asset Automation

- Implement **automate opening, closing and compliance checklists** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- task completion on time
- labor hours per unit of work
- asset downtime

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: task engine, vision/IoT signals, workforce rules, asset alerts and associate copilot; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 10 / PRACTICE 08

Build an associate copilot for procedures and exceptions


Store Operations, Workforce & Physical Asset Automation

- Implement **build an associate copilot for procedures and exceptions** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from tasks, labor, shelf state, IoT telemetry, asset health, SOPs, schedules and store events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- task completion on time
- labor hours per unit of work
- asset downtime

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: task engine, vision/IoT signals, workforce rules, asset alerts and associate copilot; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 11

Food Safety, Traceability, Recall & Compliance

Design safety and traceability as data products with rapid containment, evidence preservation and regulatory readiness.

What good looks like

- One accountable business owner owns the decision and its KPI; technology ownership is explicit but does not replace operational accountability.
- Inputs are reconciled to systems of record with freshness and quality SLOs; recommendations carry confidence, timestamp, reason codes and version lineage.
- Automation is progressive: observe -> recommend -> approve -> auto-execute low-risk cases. High-impact exceptions retain human control and audit evidence.
- The capability is evaluated on business outcomes and operational reliability, not model accuracy alone. A fallback mode keeps stores and supply chain running.

Best practices in this chapter

11.01 Map critical tracking events and key data elements

11.05 Monitor HACCP and critical control point exceptions

11.02 Implement interoperable lot-level traceability

11.06 Build a recall impact graph and containment workflow

11.03 Use GS1 identifiers and EPCIS-style event records

11.07 Preserve audit evidence and chain of custody

11.04 Automate supplier food-safety document controls

11.08 Test recall readiness with simulation exercises

- Authoritative inputs: lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events.
- Expected implementation assets: traceability store, event standards, recall graph, evidence vault and test protocol.
- Required cross-cutting controls: identity, audit, data quality, observability, rollback and human escalation.

MINIMUM CAPABILITY FOUNDATION

CHAPTER 11 / PRACTICE 01

Map critical tracking events and key data elements


Food Safety, Traceability, Recall & Compliance

- Implement **map critical tracking events and key data elements** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- trace completeness
- time to identify affected lots
- recall simulation pass rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: traceability store, event standards, recall graph, evidence vault and test protocol; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 11 / PRACTICE 02

Implement interoperable lot-level traceability


Food Safety, Traceability, Recall & Compliance

- Implement **implement interoperable lot-level traceability** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use globally unique product/location/lot identifiers, append-only event history and parent-child lot links; make backward and forward trace queries fast and testable. Use lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- trace completeness
- time to identify affected lots
- recall simulation pass rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: traceability store, event standards, recall graph, evidence vault and test protocol; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 11 / PRACTICE 03

Use GS1 identifiers and EPCIS-style event records


Food Safety, Traceability, Recall & Compliance

- Implement **use gs1 identifiers and epcis-style event records** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Map internal identifiers to standards such as GTIN/GLN and use EPCIS-style event semantics where appropriate so partners can exchange traceability events consistently. Use lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- trace completeness
- time to identify affected lots
- recall simulation pass rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: traceability store, event standards, recall graph, evidence vault and test protocol; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 11 / PRACTICE 04

Automate supplier food-safety document controls

Food Safety, Traceability, Recall & Compliance

- Implement **automate supplier food-safety document controls** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- trace completeness
- time to identify affected lots
- recall simulation pass rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: traceability store, event standards, recall graph, evidence vault and test protocol; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 11 / PRACTICE 05

Monitor HACCP and critical control point exceptions

Food Safety, Traceability, Recall & Compliance

- Implement **monitor haccp and critical control point exceptions** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Automate monitoring and evidence collection, but keep qualified food-safety personnel responsible for critical limits, corrective actions and release decisions. Use lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- trace completeness
- time to identify affected lots
- recall simulation pass rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: traceability store, event standards, recall graph, evidence vault and test protocol; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 11 / PRACTICE 06

Build a recall impact graph and containment workflow


Food Safety, Traceability, Recall & Compliance

- Implement **build a recall impact graph and containment workflow** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Build a graph from supplier lot through transformation, DC, store/order and customer exposure; support immediate quarantine, stop-sale, notification and evidence capture. Use lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- trace completeness
- time to identify affected lots
- recall simulation pass rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: traceability store, event standards, recall graph, evidence vault and test protocol; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 11 / PRACTICE 07

Preserve audit evidence and chain of custody

Food Safety, Traceability, Recall & Compliance

- Implement **preserve audit evidence and chain of custody** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- trace completeness
- time to identify affected lots
- recall simulation pass rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: traceability store, event standards, recall graph, evidence vault and test protocol; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 11 / PRACTICE 08

Test recall readiness with simulation exercises


Food Safety, Traceability, Recall & Compliance

- Implement **test recall readiness with simulation exercises** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Build a graph from supplier lot through transformation, DC, store/order and customer exposure; support immediate quarantine, stop-sale, notification and evidence capture. Use lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from lot genealogy, CTE/KDE records, supplier certificates, temperatures, inspections and recall events. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- trace completeness
- time to identify affected lots
- recall simulation pass rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: traceability store, event standards, recall graph, evidence vault and test protocol; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 12

AI Agents, Workflow Automation & Integration Architecture

Use agents as controlled decision workers inside deterministic business workflows, not as unbounded autonomous systems.
What good looks like

- One accountable business owner owns the decision and its KPI; technology ownership is explicit but does not replace operational accountability.
- Inputs are reconciled to systems of record with freshness and quality SLOs; recommendations carry confidence, timestamp, reason codes and version lineage.
- Automation is progressive: observe -> recommend -> approve -> auto-execute low-risk cases. High-impact exceptions retain human control and audit evidence.
- The capability is evaluated on business outcomes and operational reliability, not model accuracy alone. A fallback mode keeps stores and supply chain running.

Best practices in this chapter

12.01 Choose workflow automation versus AI agent patterns

12.05 Add approval gates for financial and operational actions

12.02 Design tool contracts for agent actions

12.06 Implement retries, idempotency and dead-letter handling

12.03 Use RAG for policy, SOP and product knowledge

12.07 Use n8n-style orchestration for cross-system workflows

12.04 Orchestrate multi-agent work with a supervisor pattern

12.08 Build agent evaluation, tracing and prompt lifecycle management

- Authoritative inputs: APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals.
- Expected implementation assets: workflow engine, agent runtime, RAG service, tool gateway, human gates and observability.
- Required cross-cutting controls: identity, audit, data quality, observability, rollback and human escalation.

MINIMUM CAPABILITY FOUNDATION

CHAPTER 12 / PRACTICE 01

Choose workflow automation versus AI agent patterns


AI Agents, Workflow Automation & Integration Architecture

- Implement **choose workflow automation versus ai agent patterns** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Prefer deterministic workflows for well-defined steps and rules; invoke an LLM only for unstructured interpretation, ranking or generation where errors are contained. Use APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- automated completion rate
- agent task success
- human escalation / policy violation rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: workflow engine, agent runtime, RAG service, tool gateway, human gates and observability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 12 / PRACTICE 02

Design tool contracts for agent actions


AI Agents, Workflow Automation & Integration Architecture

- Implement **design tool contracts for agent actions** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Expose narrow typed tools with schema validation, authorization checks, idempotency keys, rate limits and auditable side effects; never give agents broad database credentials. Use APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- automated completion rate
- agent task success
- human escalation / policy violation rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: workflow engine, agent runtime, RAG service, tool gateway, human gates and observability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 12 / PRACTICE 03

Use RAG for policy, SOP and product knowledge


AI Agents, Workflow Automation & Integration Architecture

- Implement **use rag for policy, sop and product knowledge** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Chunk by business meaning, attach permissions and freshness metadata, retrieve before generation, cite the source internally, and evaluate retrieval separately from answer quality. Use APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- automated completion rate
- agent task success
- human escalation / policy violation rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: workflow engine, agent runtime, RAG service, tool gateway, human gates and observability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 12 / PRACTICE 04

Orchestrate multi-agent work with a supervisor pattern


AI Agents, Workflow Automation & Integration Architecture

- Implement **orchestrate multi-agent work with a supervisor pattern** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use a supervisor to decompose work and specialized agents with limited tools; cap turns, enforce budgets and avoid agent-to-agent loops without deterministic termination. Use APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- automated completion rate
- agent task success
- human escalation / policy violation rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: workflow engine, agent runtime, RAG service, tool gateway, human gates and observability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 12 / PRACTICE 05

Add approval gates for financial and operational actions


AI Agents, Workflow Automation & Integration Architecture

- Implement **add approval gates for financial and operational actions** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define thresholds for money, safety, inventory write-offs, supplier changes and customer compensation; require explicit approval tokens for irreversible actions. Use APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- automated completion rate
- agent task success
- human escalation / policy violation rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: workflow engine, agent runtime, RAG service, tool gateway, human gates and observability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 12 / PRACTICE 06

Implement retries, idempotency and dead-letter handling


AI Agents, Workflow Automation & Integration Architecture

- Implement **implement retries, idempotency and dead-letter handling** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Assign workflow and business-event IDs, make writes idempotent, use exponential backoff, dead-letter queues and compensating actions for partial failure. Use APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- automated completion rate
- agent task success
- human escalation / policy violation rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: workflow engine, agent runtime, RAG service, tool gateway, human gates and observability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 12 / PRACTICE 07

Use n8n-style orchestration for cross-system workflows


AI Agents, Workflow Automation & Integration Architecture

- Implement **use n8n-style orchestration for cross-system workflows** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Separate credentials from workflows, use sub-workflows for reusable controls, validate inputs, persist state externally for critical jobs and route failures to an owned exception queue. Use APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- automated completion rate
- agent task success
- human escalation / policy violation rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: workflow engine, agent runtime, RAG service, tool gateway, human gates and observability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 12 / PRACTICE 08

Build agent evaluation, tracing and prompt lifecycle management


AI Agents, Workflow Automation & Integration Architecture

- Implement **build agent evaluation, tracing and prompt lifecycle management** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Maintain golden datasets and production traces, test task success, groundedness, tool choice and policy compliance, and block releases that regress critical metrics. Use APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from APIs, events, knowledge bases, tool permissions, workflow state, agent traces and approvals. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- automated completion rate
- agent task success
- human escalation / policy violation rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: workflow engine, agent runtime, RAG service, tool gateway, human gates and observability; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 13

Platform Engineering, Security, Governance & FinOps

Operate AI and automation as a governed production platform with security, observability, cost control and reliable change management.

What good looks like

- One accountable business owner owns the decision and its KPI; technology ownership is explicit but does not replace operational accountability.
- Inputs are reconciled to systems of record with freshness and quality SLOs; recommendations carry confidence, timestamp, reason codes and version lineage.
- Automation is progressive: observe -> recommend -> approve -> auto-execute low-risk cases. High-impact exceptions retain human control and audit evidence.
- The capability is evaluated on business outcomes and operational reliability, not model accuracy alone. A fallback mode keeps stores and supply chain running.

Best practices in this chapter

13.01 Create an AI platform reference architecture

13.05 Monitor latency, errors, quality and business outcomes

13.02 Apply least privilege, secrets management and network controls

13.06 Control LLM and automation cost with budgets and routing

13.03 Classify AI use cases by risk and required controls

13.07 Design disaster recovery and graceful degradation

13.04 Implement model, prompt and workflow versioning

13.08 Run production release, rollback and incident management

- Authoritative inputs: identity, secrets, logs, model usage, workflow versions, incidents, costs and service health.
- Expected implementation assets: platform controls, policy-as-code, observability, cost governance, release process and DR plan.
- Required cross-cutting controls: identity, audit, data quality, observability, rollback and human escalation.

MINIMUM CAPABILITY FOUNDATION

CHAPTER 13 / PRACTICE 01

Create an AI platform reference architecture

Platform Engineering, Security, Governance & FinOps

- Implement **create an ai platform reference architecture** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define the decision boundary explicitly, use authoritative system data, make the recommendation explainable, keep irreversible actions behind policy controls, and measure the closed-loop business outcome. Use identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1	2	3	4	5
Baseline & owner	Trusted data	Decision logic	Shadow mode	Progressive automation

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- platform availability
- cost per successful task
- change failure / rollback rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: platform controls, policy-as-code, observability, cost governance, release process and DR plan; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 13 / PRACTICE 02

Apply least privilege, secrets management and network controls


Platform Engineering, Security, Governance & FinOps

- Implement **apply least privilege, secrets management and network controls** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use workload identities, short-lived credentials, scoped roles, private networking where practical and separate read from write tools for AI agents. Use identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- platform availability
- cost per successful task
- change failure / rollback rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: platform controls, policy-as-code, observability, cost governance, release process and DR plan; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 13 / PRACTICE 03

Classify AI use cases by risk and required controls


Platform Engineering, Security, Governance & FinOps

- Implement **classify ai use cases by risk and required controls** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Classify use cases by impact, autonomy, data sensitivity and reversibility; require stronger evidence and human control as consequence severity rises. Use identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- platform availability
- cost per successful task
- change failure / rollback rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: platform controls, policy-as-code, observability, cost governance, release process and DR plan; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 13 / PRACTICE 04

Implement model, prompt and workflow versioning


Platform Engineering, Security, Governance & FinOps

- Implement **implement model, prompt and workflow versioning** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Version model, prompt, tools, retrieval index, workflow and policy together; record the exact release used for every important decision. Use identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- platform availability
- cost per successful task
- change failure / rollback rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: platform controls, policy-as-code, observability, cost governance, release process and DR plan; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 13 / PRACTICE 05

Monitor latency, errors, quality and business outcomes

Platform Engineering, Security, Governance & FinOps

- Implement **monitor latency, errors, quality and business outcomes** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Set service-level objectives for end-to-end business tasks, not only model calls; trace dependency time, queueing, retries and human wait states. Use identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- platform availability
- cost per successful task
- change failure / rollback rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: platform controls, policy-as-code, observability, cost governance, release process and DR plan; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 13 / PRACTICE 06

Control LLM and automation cost with budgets and routing

Platform Engineering, Security, Governance & FinOps

- Implement **control llm and automation cost with budgets and routing** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Track cost per completed business outcome, route simple tasks to smaller models, cache stable results, cap context and set per-workflow budgets and alerts. Use identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- platform availability
- cost per successful task
- change failure / rollback rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: platform controls, policy-as-code, observability, cost governance, release process and DR plan; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 13 / PRACTICE 07

Design disaster recovery and graceful degradation


Platform Engineering, Security, Governance & FinOps

- Implement **design disaster recovery and graceful degradation** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Define degraded modes that preserve sales and safety when AI is unavailable; maintain deterministic fallbacks, queued work and tested recovery procedures. Use identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN



- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- platform availability
- cost per successful task
- change failure / rollback rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: platform controls, policy-as-code, observability, cost governance, release process and DR plan; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

CHAPTER 13 / PRACTICE 08

Run production release, rollback and incident management

Platform Engineering, Security, Governance & FinOps

- Implement **run production release, rollback and incident management** as a measurable decision capability, connected to operating systems and owned by the business team that acts on the result.

TARGET OUTCOME

- Use canary releases, feature flags, rollback, runbooks and post-incident reviews; preserve traces and business-event IDs for root-cause analysis. Use identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. The production design should expose confidence, reason codes, source timestamps and the policy version used.

RECOMMENDED PRODUCTION DESIGN

1

Baseline & owner

- Baseline the current process, decision latency, manual effort, error modes and business KPI. Name an accountable product owner and operational approver.
- Create the minimum trusted data product from identity, secrets, logs, model usage, workflow versions, incidents, costs and service health. Validate freshness, completeness, entity keys and reconciliation against the system of record.
- Implement the decision logic in stages: deterministic rules first, optimization/ML second, and LLM interpretation only where unstructured information is required.
- Run shadow mode against live work. Compare recommendations with actual decisions, capture overrides/reasons, test failure paths and tune thresholds before automation.
- Progressively automate low-risk cases. Monitor business outcomes, drift, cost and exceptions; retain rollback, manual fallback and an auditable release record.

2

Trusted data

3

Decision logic

4

Shadow mode

5

Progressive automation

- No silent writes: every side effect carries actor/workflow ID, reason, timestamp and idempotency key.
- Human approval for high-impact, irreversible, safety-sensitive, employment, supplier master, bank/payment or large financial actions.
- Fail closed when required data is stale/invalid; fail gracefully to deterministic fallback when AI/model services are unavailable.
- Separate recommendation quality from execution reliability; alert the owning team on both model and workflow failures.

IMPLEMENTATION SEQUENCE

- platform availability
- cost per successful task
- change failure / rollback rate

KPIS / ACCEPTANCE GATES

CONTROLS & FAILURE HANDLING

- Production-ready v1 package: platform controls, policy-as-code, observability, cost governance, release process and DR plan; plus acceptance tests, dashboard, SOP/runbook, owner/RACI and a 30-day post-launch measurement plan.

FIRST IMPLEMENTATION DELIVERABLE

Canonical Retail Data Model

Use stable business identifiers and event history so ML, optimization and agents do not invent their own version of reality.

Entity	Required attributes / design notes
Product	GTIN/SKU, description, hierarchy, dimensions, pack, allergens, temperature class, shelf-life rules
Location	GLN/internal ID, store/DC/supplier type, timezone, capacity, region, fulfillment capabilities
InventoryPosition	product-location, on-hand, reserved, available, in-transit, quarantine, damaged, timestamp, confidence
Lot/Batch	lot ID, product, supplier, manufacture/pack/expiry, receipt, quality status, parent/child lineage
Order	channel, lines, promised service, node, substitutions, payment state, customer consent references
SupplyEvent	PO, ASN, receipt, transfer, shipment, delivery, adjustment, reason, actor, event time
Commercial	price, promotion, cost, assortment, planogram, competitor observation, effective dates
TraceEvent	event type, what/when/where/why, business step, disposition, source, linked lots/shipments

Example Retail Event Contract

Illustrative JSON-style schema. In production, validate schema versions and business invariants at ingestion and before action.

```
{
  "event_id": "01J...",
  "event_type": "inventory.receipt.completed",
  "schema_version": "2.1",
  "occurred_at": "2026-08-19T14:22:31-06:00",
  "source": {"system": "WMS", "location_id": "DC-104"},
  "subject": {"product_id": "SKU-8812", "lot_id": "LOT-44A"},
  "quantity": {"value": 120, "uom": "EA"},
  "attributes": {
    "expiry_date": "2026-08-27",
    "temperature_c": 3.2,
    "quality_status": "RELEASED"
  },
  "trace": {"po_id": "PO-90017", "asn_id": "ASN-2218"},
  "actor": {"type": "service", "id": "receiving-v3"},
  "idempotency_key": "DC-104|ASN-2218|SKU-8812|LOT-44A",
  "data_classification": "INTERNAL"
}
```

- event_id must be globally unique; event times include timezone; product/location/lot keys resolve to mastered entities.
- Quantity and unit-of-measure must pass conversion rules; expiry/temperature fields are required for configured perishable categories.
- Consumers must be idempotent. Schema-breaking changes create a new version; old consumers receive a deprecation window.
- Store provenance, actor, policy/schema version and trace IDs so every downstream automated action can be reconstructed.

VALIDATION RULES

Replenishment Decision Pseudocode

Keep mathematical/constraint logic deterministic and let AI assist with exception interpretation, not unrestricted ordering.


```
for sku_location in eligible_items:
    inv = trusted_inventory_position(sku_location)
    if inv.confidence < MIN_CONFIDENCE: route_to_cycle_count(); continue
    demand = forecast_quantiles(sku_location, horizon=lead_time + review_period)
    target = demand.P50 + safety_stock(demand, lead_time_uncertainty, service_target)
    net_need = target - inv.available - confirmed_inbound + backorders
    qty = round_to_pack(max(0, net_need), case_pack)
    qty = apply_constraints(qty, shelf_life, MOQ, capacity, supplier_limits, budget)
    if exception_risk(qty) <= AUTO_RELEASE_THRESHOLD:
        create_po_line(idempotency_key=decision_id)
    else:
        queue_for_approval(reason_codes, evidence, expected_service_impact)
```

- Backtest service level, waste and inventory investment against the current policy; do not approve solely on forecast error.
- Replay historical stockouts, late suppliers, pack-size changes and bad inventory records; verify the engine refuses unsafe recommendations.
- Log each input snapshot, output, constraint binding and override. Use shadow mode before auto-release.
- Cap daily/weekly order-value change during rollout and maintain emergency manual ordering.

PRODUCTION ACCEPTANCE

Governed Retail Agent Pattern

Agents operate inside a workflow envelope with narrow tools, explicit policy checks and human approval for consequential actions.



- retrieve SOP/policy/product facts within authorization
- classify exceptions and propose ranked actions
- call read-only tools and low-risk write tools inside thresholds
- draft customer/supplier/associate communications for approved workflows

AGENT MAY DO

- change supplier bank details, approve payments or large POs without a deterministic approval service
- release quarantined food or override safety holds
- make high-impact employee decisions from opaque model outputs
- write directly to core databases using broad credentials

AGENT MUST NOT DO

- prompt/model/tool/retrieval versions
- tool arguments and results with sensitive fields redacted
- policy decision and approval identity
- task success, latency, cost, retries, escalation and user override

REQUIRED TELEMETRY

Retail AI KPI Dictionary

Every automation should have a business outcome metric, quality metric, operational reliability metric and safety/risk guardrail.

Area	Primary business metrics	Quality / guardrails
Availability	On-shelf availability, fill rate, stockout rate, ATP promise accuracy	Inventory confidence, substitution rate
Inventory	Turns, days of supply, working capital, excess/obsolete	Reconciliation error, forecast bias
Fresh / Waste	Shrink %, waste value, sell-through before expiry	Expiry capture, cold-chain excursion response
Supply	OTIF, lead-time variability, inbound fill, cost per case	ASN match, supplier exception cycle time
Commercial	Incremental gross profit, promo ROI, margin	Price guardrail violations, elasticity stability
Fulfillment	Cost/order, pick rate, on-time delivery, perfect order	Routing fallback rate, promise error
Automation	Touchless completion, cycle time, labor hours saved	Retry/dead-letter rate, human override
AI / Agent	Task success, groundedness, tool success, escalation	Policy violations, unsafe actions, hallucination rate
Platform	Availability, p95 latency, change failure rate, MTTR	Cost/successful task, budget variance

12-Month Implementation Roadmap

Sequence foundational data and low-risk automation before scaling autonomous decisions.

0-60 days	Foundation Value-stream baseline; product/location/supplier master gaps; inventory reconciliation; event/API inventory; governance; first KPI dashboards.
2-4 months	Pilot Forecast/replenishment shadow mode; supplier/invoice automation; SOP RAG copilot; one store/DC exception workflow; traceability event prototype.
4-7 months	Controlled automation Auto-release low-risk replenishment; fresh markdown pilot; ATP/order-routing improvements; tool-gated agents; observability and evaluation platform.
7-10 months	Scale Expand to categories/locations; DC/transport optimization; store task orchestration; supplier risk; recall graph; FinOps and SRE operating model.
10-12 months	Optimize Causal promo/pricing, network scenarios, multi-agent exception handling, continuous experimentation, benefits realization and portfolio reprioritization.

Team & RACI Blueprint

Keep domain accountability with retail operators while platform teams provide reusable capabilities.

- Executive sponsor: owns strategic outcome and funding.
- Domain product owner: owns KPI, policy, backlog and rollout.
- Process SMEs: validate edge cases, safety constraints and operational fit.
- Store/DC/supplier change lead: training, SOP, adoption and feedback.

BUSINESS PRODUCT TEAM

- Data product owner + engineers: entity/event quality and contracts.
- ML/optimization engineers: forecasting/decision services and monitoring.
- Automation/agent engineers: workflows, tools, approvals, RAG and evals.
- Platform/SRE/security: identity, observability, release, resilience, incident and cost controls.

TECHNOLOGY TEAM

- One named person is Accountable for each business KPI. AI or an IT team is never the accountable owner. High-impact policy overrides have a named human approver and escalation path.

DECISION RULE

Vendor / Platform Evaluation Scorecard

Score platforms against your retail architecture and operating requirements, not demos.

- Integration/API/event support - 20%
- Reliability, scaling, observability and support - 15%
- Security, identity, data residency and audit - 20%
- AI/agent capability, evaluation and tool control - 15%
- Workflow state, retries, idempotency and human approval - 10%
- Total cost, portability and exit strategy - 10%
- Retail fit, ecosystem and team skill alignment - 10%

WEIGHTED CRITERIA

- Use one real workflow with real edge cases; measure end-to-end task success, p95 latency, failure recovery, operator effort and cost.
- Require evidence for role-based access, secrets rotation, audit logs, version rollback, data deletion and production support.
- Test a vendor outage and model outage before selection; verify exportability of workflow definitions, data and traces.

PROOF-OF-VALUE TEST

AI / Automation Risk Classification

Scale controls with consequence, autonomy, data sensitivity and reversibility.

- Draft/explain/summarize with no direct write; approved internal data only. Controls: authentication, RAG permissions, logs, user review.

TIER 1 - ASSIST

- Ranking/forecast/decision recommendation affects work but human confirms. Controls: quality gates, reasons/confidence, overrides, drift and audit.

TIER 2 - RECOMMEND

- Automated low-risk writes inside deterministic thresholds. Controls: narrow tools, idempotency, policy engine, rate/value caps, rollback and continuous monitoring.

TIER 3 - BOUNDED ACTION

- Safety, large financial, employee, sensitive customer or irreversible decisions. Default: qualified human approval and domain-specific governance; autonomous action only with exceptional evidence and controls.

TIER 4 - HIGH IMPACT

Production Go-Live Checklist

A minimum release gate for any AI-enabled retail automation that can affect operations, customers, suppliers or money.

- Owner/RACI signed; decision boundary documented; current baseline captured; target KPI and stop criteria agreed.
- SOP updated; exception ownership staffed; training completed; manual fallback tested.

BUSINESS & PROCESS

- Entity keys reconcile; freshness/completeness SLOs active; historical backfill verified; privacy/classification and retention applied.
- Bad/stale/missing input tests pass; lineage from source to decision is queryable.

DATA

- Baseline comparison passes; representative and stress datasets pass; uncertainty/thresholds calibrated; drift monitors configured.
- Reason codes and explanation are understandable to operators; known limitations documented.

MODEL / DECISION

- Idempotency, retry, timeout, dead-letter and compensation tested; duplicate-event replay safe.
- Human approval gates enforced server-side; irreversible action tests pass; rollback/kill switch validated.

WORKFLOW

- Least-privilege identities; secrets rotated; tool scopes tested; prompt injection/data exfiltration scenarios tested for agents.
- Food-safety, financial, employee and customer-impact constraints reviewed by accountable domain owners.

SECURITY & SAFETY

- Dashboards/alerts on business, model and workflow KPIs; runbook/on-call owner; incident severity matrix; capacity and cost budgets.
- Canary rollout plan, change record, release versions and 30-day post-launch review scheduled.

OPERATIONS

Incident & Rollback Playbook

Design incident response around business impact and affected decisions, not only technical errors.

- Alert on stockout/waste/service anomalies, agent policy violations, duplicate writes, unusual order/payment volume, stale data, drift, queue backlog and dependency failure.

DETECT

- Disable feature flag or write tool; switch to deterministic fallback; quarantine suspect recommendations; cap automated volumes; preserve evidence and trace IDs.

CONTAIN

- Rollback model/prompt/workflow/policy bundle; replay safe idempotent events; reconcile systems of record; verify downstream customer/supplier impact.

RECOVER

- Root-cause by data, model, workflow, policy, integration or human-process failure; add regression tests; update runbook, thresholds and owner training.

LEARN

Food Traceability Implementation Note

Build traceability as a continuously validated operational capability, not a compliance spreadsheet.

- Define product/lot/location identifiers; capture receiving, transformation and shipping events; preserve parent-child lot genealogy and source timestamps.
- Map required Critical Tracking Events / Key Data Elements for covered foods and trading-partner exchanges; retain records according to applicable requirements.

DATA DESIGN

- Reject or exception-route incomplete trace events at receiving/shipping. Connect quarantine and stop-sale to affected lot/location records.
- Build backward/forward queries and customer/order exposure where lawful; test response under time pressure.

OPERATIONAL WORKFLOW

- The FDA Food Traceability Rule has had compliance/enforcement timing changes. Validate the current FDA page and legal advice before setting a production compliance deadline; do not rely on a handbook publication date alone.

2026 PLANNING NOTE

30-60-90 Day Starter Plan

A realistic way to start without trying to automate the entire retailer.

- Choose 2 value streams: one supply/inventory and one back-office/service workflow. Baseline KPIs and exception volumes.
- Inventory systems/APIs; reconcile master data; define event IDs and audit requirements; set risk tiers and approval thresholds.
- Build a small cross-functional squad and production definition-of-done.

DAYS 0-30

- Launch data quality dashboard and first canonical events. Run forecast/replenishment or invoice workflow in shadow mode.
- Deploy one policy/SOP RAG assistant with read-only tools and evaluation set.
- Create operational dashboards, dead-letter queue and on-call/runbook before write automation.

DAYS 31-60

- Auto-execute only low-risk cases that pass acceptance thresholds; keep daily caps and kill switch.
- Measure real availability/waste/cycle-time value; collect override reasons; expand only if value and reliability hold.
- Prioritize next quarter from evidence, not vendor feature lists.

DAYS 61-90

Current Public References & Further Reading

Use these sources to validate platform capabilities, regulatory requirements and standards during implementation.

Walmart corporate / Global Tech

Public material on AI-enabled supply chain, real-time inventory decisions, route optimization, associate tools, defect detection, forecasting and retail automation.

<https://corporate.walmart.com/> ; <https://tech.walmart.com/>

U.S. FDA - Food Traceability Rule

Rule, compliance guidance, current implementation/enforcement updates, Critical Tracking Events and Key Data Elements.

<https://www.fda.gov/food/food-safety-modernization-act-fsma/fsma-final-rule-requirements-additional-traceability-records-certain-foods>

GS1

GTIN/GLN identifiers, retail/grocery standards, EPCIS and supply-chain visibility / traceability guidance.

<https://www.gs1.org/industries/retail> ; <https://www.gs1.org/standards/epcis>

NIST AI Risk Management Framework

AI RMF and Generative AI Profile for governance, measurement and risk treatment.

<https://www.nist.gov/itl/ai-risk-management-framework>

OpenAI developer resources

Current agent, tool-use, evaluation, security and API guidance.

<https://platform.openai.com/docs/> ; <https://openai.github.io/openai-agents-python/>

n8n documentation

Workflow orchestration, credentials, error handling, queue/scaling and AI workflow patterns.

<https://docs.n8n.io/>